


## Gather/Scatter On Pytorch

Those operations are used for selecting/distributing values on a tensor, but in a fast way. For example implementing Cross-Entropy using gather has good performance, unfortunately the documentation of those functions are not very explanatory.


### Gather

Get values from input tensor along a dimension dim using some indexes. Both input and index needs to have the same shape.


#### Gather in 1D



Gathering along the dimension 1, means that we're addressing the input like this:



Doing the same on the dimension zero would have a completely different effect, on this case we would not even be able to address the input properly because there is no index besides 0



An important point to observe is that the index reset to 0 at the end of the column. And yes it's confusing but dimension 0 scan columnwise while dimension 1 scan rowise.

```

a = torch.tensor([[1,2,3]])
b = torch.tensor([[1,0,2]])
r = torch.gather(input=a, dim=0, index=b)
print(r)

```

RuntimeError Traceback (most recent call last)

```

<ipython-input-77-cdf908c75462> in <module>
      1 a = torch.tensor([[1,2,3]])
      2 b = torch.tensor([[1,0,2]])
----> 3 r = torch.gather(input=a, dim=0, index=b)
      4 print(r)

```

RuntimeError: Invalid index in gather at /opt/conda/conda-bld/pytorch\_1556653215914/work/aten/src/TH/generic/THTensorEvenMoreMath.cpp:459

The only way this could work would be to use zeros everywhere on the index.

```

a = torch.tensor([[1,2,3]])
b = torch.tensor([[0,0,0]])
r = torch.gather(input=a, dim=0, index=b)
print(r)

```

tensor([[1, 2, 3]])

## Gather in 2D


In 2D things get more interesting, here you can consider the index as a placeholder with same shape as input where every cell on the index it's a pointer to where to gather on the input, consider the following example:

```


input = torch.tensor([[1,2],[3,4]])
index = torch.tensor([[0,0],[1,0]])
r = torch.gather(input, 0, index)
print(r)

```


tensor([[1, 2],  
 [3, 2]])




Here gathering along the dimension zero means that we're addressing the input like this:



So if we put together the index matrix we would have this:



Now consider that we're choose the dimension 1



In practice people reshape the inputs to a vector to make things more manageable in their heads.

### **Scatter**

This operation is used to scatter values along a tensor following the same rules as gather, but now pushing values instead of selecting them.

```
x = torch.arange(3).view(1,3)+10
print('input:',x)
y = torch.zeros_like(x)
# index need to have same shape as x
b = torch.tensor([[0,2,1]])
print('index:',b)
# Scatter on y, the values on x along dimension 1
y.scatter(dim=1, index=b, src=x)
```

```
input: tensor([[10, 11, 12]])
index: tensor([[0, 2, 1]])
tensor([[10, 12, 11]])
```

## Why this?

- Because it's faster than doing in a for-loop
- Pytorch uses same semantic for distributing tensors across different machines/GPUs

## Example of CrossEntropy

Here we implement a CrossEntropy with gather, one cool feature is that gather accept broadcast so the input could be [batch x classes] while the index [batch x 1]

```
# Define custom cross_entropy
# x: shape [batch x C]
# y: shape [batch]
def my_cross_entropy(x, target):
    # Calculate the log-probability of x along first dimension
    log_prob = -1.0 * F.log_softmax(x, 1)
    print('log_prob:')
    print(log_prob)
    # Unsqueeze will make y shape become [batch x 1]
    # Gather elements from log_prob along second dimension (dim=1)
    # here log_prob.shape [ batch x C]
    # target.shape [ batch x 1], so the shapes wont match but pytorch will broadcast
    loss_no_gather = log_prob[range(log_prob.shape[0]), target]
    print('No gather loss(SLOW):')
    print(loss_no_gather.unsqueeze(1))
    loss = log_prob.gather(1, target.unsqueeze(1))
    print('Gather loss')
    print(loss)
    loss = loss.mean()
    return loss

# Reference CrossEntropy
criterion = nn.CrossEntropyLoss()
```

Inspecting closer what's going on you can observe:

- Labels after `unsqueeze(1)` will become [batch x 1] which doesn't exactly match with [batch x C]

- Gather will broadcast and select the correct index for each element on the batch along the dimension 1
- Both custom Cross\_Entropy and Reference Cross entropy return the same value
- The other option would be `loss_no_gather = log_prob[range(log_prob.shape[0]), target]` (SLOW because doesn't use pytorch ops)

```

labels:
tensor([[3],
        [2],
        [3]])
      Labels will act as indexes
predictions:
tensor([[ -0.9637, -1.6281,  0.4499, -0.3091,  1.5096],
        [ 0.6513,  1.1302,  1.2812, -0.3062, -1.6226],
        [ 0.5829, -0.6996,  1.0261,  0.0362,  0.3030]], requires_grad=True)
x.shape: torch.Size([3, 5])
y.shape: torch.Size([3])
log_prob:
tensor([[2.9659, 3.6303, 1.5523, 2.3113, 0.4925],
        [1.6051, 1.1263, 0.9752, 2.5626, 3.8790],
        [1.4278, 2.7103, 0.9846, 1.9746, 1.7078]])
      Stuff gathered from log_prob
Gather loss
tensor([[2.3113],
        [0.9752],
        [1.9746]])
loss_reference: tensor(1.7537)
my_cross_entropy: tensor(1.7537)

```

## References

- <https://discuss.pytorch.org/t/convert-int-into-one-hot-format/507/3>
- <https://discuss.pytorch.org/t/custom-loss-functions/29387/8>
- <https://stackoverflow.com/questions/50999977/what-does-the-gather-function-do-in-pytorch-in-layman-terms>